

#### What does this session cover?

What is the TempDB?

**Types of TempDB Contention** 

**Optimizing the TempDB Database** 

**TempDB Performance Improvements** 



### What is the TempDB database?

#### System database

- Available to all users with the same structure as user databases.
- Operations are minimally logged.
- Re-created every time SQL Server is started.

#### Workload

- Used for temporary (nondurable) storage.
- Object and data frequently being created and destroyed.
- Very high concurrency.
- Backup and restore operations are not allowed on TempDB.

#### What is stored in TempDB?

# Temporary user objects

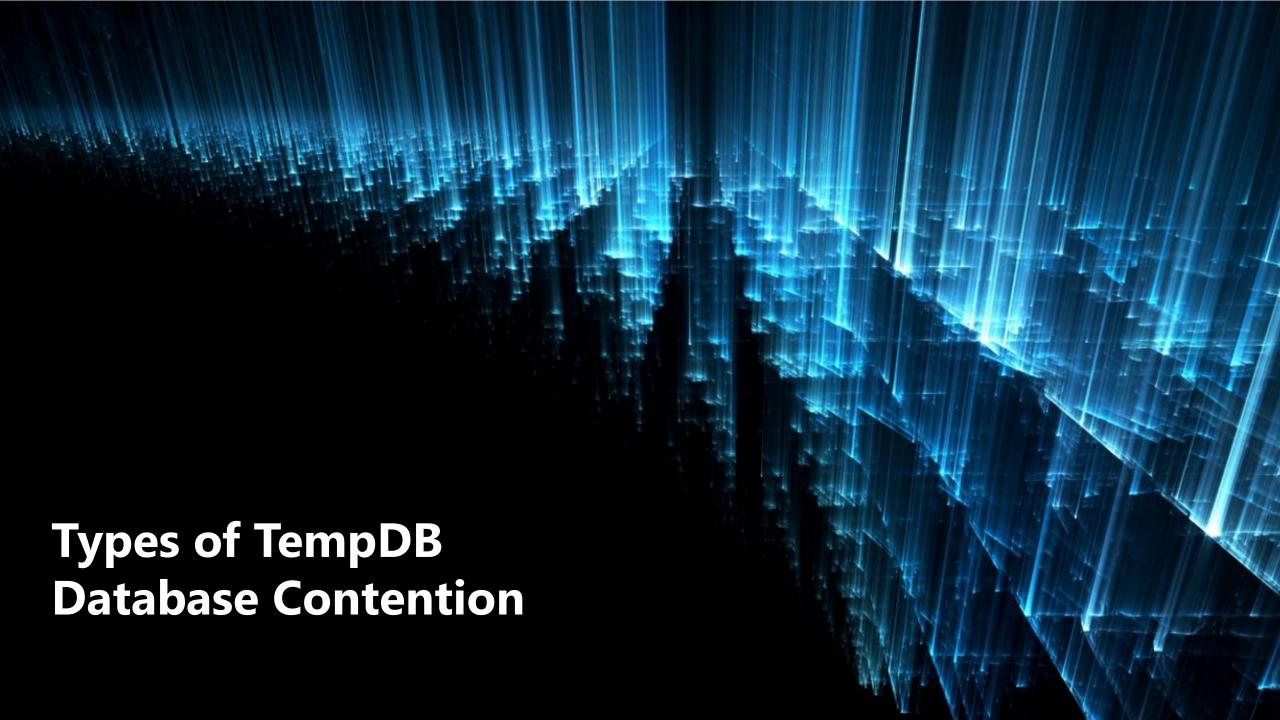
- Global or local temporary tables and indexes
- Temporary stored procedures
- Table variables
- Tables returned in table-valued functions

#### Internal objects

- Worktables to store intermediate results for spools, cursors, sorts, and temporary LOB storage.
- Work files for hash join or hash aggregate operations.

Version stores

- Common row version store and online-index-build version store
- Version stores can be moved to user databases by enabling Accelerated Database Recovery (ADR) in SQL Server 2019



### Types of TempDB Contention

#### Object allocation contention

- The wait type is PAGELATCH\_UP on pages in TempDB.
- These pages might be PFS (2:1:1) or SGAM (2:1:3) pages in TempDB.

#### Metadata contention

• The wait type is PAGELATCH\_EX on **sysschobjs** in TempDB.

#### Temp Table Cache Contention

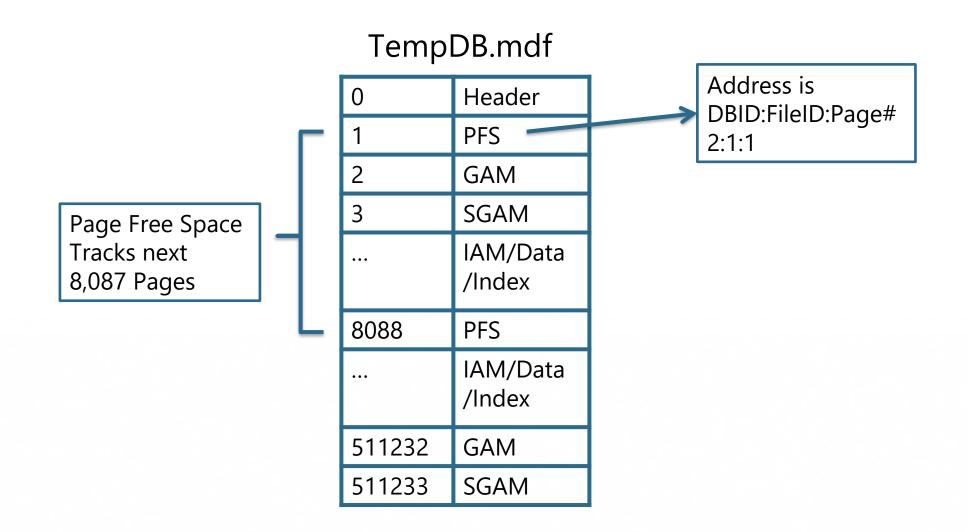
- The wait type is CMEMTHREAD or SOS\_CACHESTORE spinlock waits.
- This could indicate other issues besides TempDB Temp Table caching.

### The Roles of Allocation Pages

PFS and IAM pages are used to determine if a new page or extent is needed

GAM and SGAM pages are used to allocate extents

### Database Page Layout



#### **Metadata Contention**

#### The wait type is PAGELATCH\_EX on sysschobjs in TempDB.

session_id	wait_type	wait_resource	object_name	blocki	command	statement_text	data	file_id	page_id	object_id	index_id	page_typ
146	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
147	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
160	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
161	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
162	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
163	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
164	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
165	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
166	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
167	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
168	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
169	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
170	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
171	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
172	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
173	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F



### TempDB File Placement

File	Logical name	Physical name	Initial size	File growth
Primary data	tempdev	tempdb.mdf	8 MB	Autogrow by 64 MB until the disk is full
Secondary data files*	temp#	tempdb_mssql_#.ndf	8 MB	Autogrow by 64 MB until the disk is full
Log	templog	templog.ldf	8 MB	Autogrow by 64 megabytes to a maximum of 2 terabytes

If the number of logical processors is less than or equal to eight (8), use the same number of data files as logical processors.

If the number of logical processors is greater than eight (8), use eight data files.

If contention continues, increase the number of data files by multiples of four (4) up to the number of logical processors

Alternatively, make changes to the workload or code.

## Optimizing TempDB performance

Consider instant file initialization

Pre-allocate space for all TempDB files

Divide TempDB into multiple data files of equal size

Put the TempDB database on a fast I/O subsystem

Use disk striping if there are many directly attached disks.

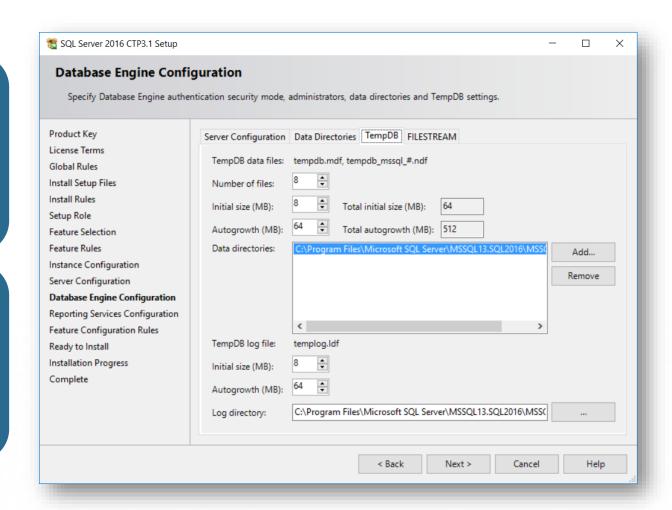
Put the TempDB database on separate disks from user databases



# Performance improvements in TempDB (2016)

Setup adds multiple TempDB data files during instance installation.

By default, setup adds as many TempDB data files as the logical processor count or eight, whichever is lower.



# Performance improvements in TempDB (2016)

Trace Flags behavior enabled by default for TempDB

- 1117 (Grow all files in a filegroup evenly)
- 1118 (Avoid mixed extents and only use full extents)

Temporary tables and table variables are cached.

Improved allocation page latching.

PAGELATCH\_SH on metadata allocation instead of PAGELATCH\_EX

Logging overhead for TempDB is reduced.

### Performance improvements in TempDB (2019)

#### Default

- Temp table cache improvements
- Concurrent PFS updates

#### Opt-in

 Memory-Optimized TempDB Metadata

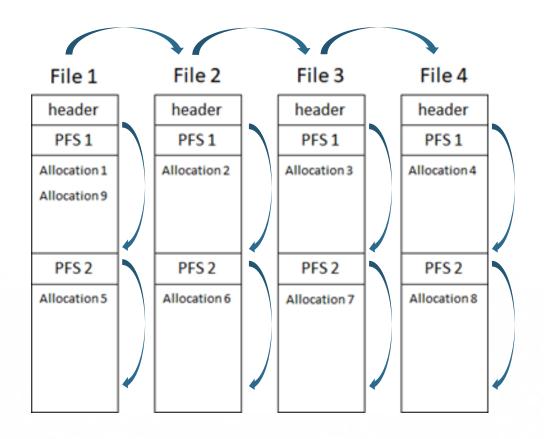
**TEMPDB Files, Trace Flags, and Updates by Pam Lahoud** 

**How (and When) To: Memory Optimized TempDB** 

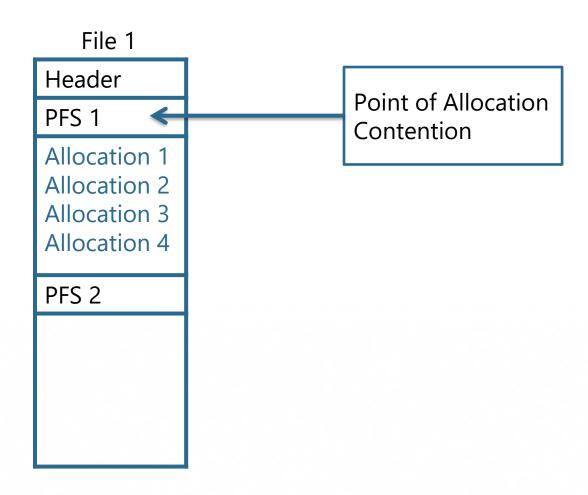
#### Concurrent PFS Pages in TempDB (2019)

New algorithm uses round robin between files, and between PFS pages within the files.

With this change, not only will increasing the number of files help with PFS contention but also increasing the size of the files.



# PFS Page Allocation (Single File: Pre-2019)



## PFS Page Round Robin (Multiple Files: Pre-2019)

File 1 File 2 File 3 File 4 Header Header Header Header PFS 1 PFS 1 PFS 1 PFS 1 Allocation 3 Allocation 1 Allocation 4 Allocation 2 Allocation 5 Allocation 6 PFS 2 PFS 2 PFS 2 PFS 2

### PFS Page Round Robin (Multiple Files: 2019)

File 1 File 2 File 3 File 4 Header Header Header Header PFS 1 PFS 1 PFS 1 PFS 1 Allocation 1 Allocation 3 Allocation 2 Allocation 4 PFS 2 PFS 2 PFS 2 PFS 2 Allocation 5 Allocation 7 Allocation 8 Allocation 6

#### **Metadata Contention**

#### The wait type is PAGELATCH\_EX on sysschobjs in TempDB.

session_id	wait_type	wait_resource	object_name	blocki	command	statement_text	data	file_id	page_id	object_id	index_id	page_typ
146	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
147	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
160	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
161	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
162	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
163	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
164	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
165	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
166	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
167	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
168	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
169	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
170	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
171	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F
172	PAGELATCH_EX	2:1:118	sysschobjs	105	CREATE TABLE	CREATE TABLE #This_Table (col1 INT)	2	1	118	34	2	INDEX_F
173	PAGELATCH_EX	2:1:118	sysschobjs	105	EXECUTE	EXEC jd_temp_demo	2	1	118	34	2	INDEX_F

### Identifying Metadata Contention

```
SELECT er.session_id, er.wait_type, er.wait_resource,
OBJECT_NAME(page_info.[object_id],page_info.[database_id]) as [object_name],
er.blocking_session_id, er.command,
SUBSTRING(st.text, (er.statement_start_offset/2) + 1,
    ((CASE statement end offset
        WHEN -1 THEN DATALENGTH(ST.text)
        ELSE er.statement_end_offset END
            - er.statement_start_offset)/2) + 1) AS statement_text,
page info.database id,
page_info.[file_id],
page_info.page_id,
page info.[object id],
page_info.index_id,
page_info.page_type_desc
FROM sys.dm exec requests AS er
  CROSS APPLY sys.dm_exec_sql_text(er.sql_handle) as st
 CROSS APPLY sys.fn_PageResCracker(er.page_resource) AS r
  CROSS APPLY sys.dm db page info(r.db id, r.file id, r.page id, 'DETAILED') AS
       page info;
```

# TempDB Memory Optimized Metadata Tables

Enable Memory Optimized TempDB Metadata Tables

ALTER SERVER CONFIGURATION SET MEMORY\_OPTIMIZED TEMPDB\_METADATA = ON;

Search for Memory Optimized Tables

SELECT OBJECT\_NAME(object\_id) as [object\_name], \* FROM sys.dm\_db\_xtp\_object\_stats;

	object_name	object_id	xtp_object_id	row_insert_attempts	row_update_attempts	row_delete_attempts
1	sysrscols	3	-2147483648	1406	0	34
2	sysseobjvalues	9	-2147483647	0	0	0
3	sysschobjs	34	-2147483644	2563	55	4
4	sysmultiobjvalues	40	-2147483643	0	0	0
5	syscolpars	41	-2147483640	1158	1	16
6	sysidxstats	54	-2147483639	210	7	7
7	sysiscols	55	-2147483638	521	0	9
8	sysobjvalues	60	-2147483637	196	0	5
9	syssingleobjrefs	74	-2147483634	208	0	4
10	sysmultiobjrefs	75	-2147483633	107	0	0

Faleminderit Shukran Chnorakaloutioun Dankie Blagodaria Hvala Tak Dank u Tänan **Merci** Danke Kiitos Ευχαριστώ Děkuji A dank Köszönöm Takk Terima kasih Mahalo תודה. Dhanyavād Grazie Grazzi

# Thank you!

ありがとうございました 감사합니다 Paldies Ačiū Choukrane Благодарам 谢谢 Obrigado Спасибо Dziękuję Multumesc Баярлалаа Ngiyabonga Kop khun Teşekkür ederim

Дякую

Хвала

Ďakujem

Tack

Nandri

Diolch